

Gradient descent for Logistic Regression

Gradient Descent for Logistic Regression

Given the cost function (J) and the gradient descent algorithms for w and b ,

$$J(\vec{w}, b) = -\frac{1}{m} \sum_{i=1}^m [-y^i \log(f_{\vec{w},b}(\vec{x}^i)) + (1 - y^i) \log(1 - f_{\vec{w},b}(\vec{x}^i))]$$

$$\text{repeat} \begin{cases} w_j = w_j - \alpha \frac{\partial}{\partial w_j} J(\vec{w}, b) \\ b = b - \alpha \frac{\partial}{\partial b} J(\vec{w}, b) \end{cases}$$

⇒ Simultaneous update

The derivative can be calculated as:

$$\begin{aligned} \frac{\partial}{\partial w} J(\vec{w}, b) &= \frac{1}{m} \sum_{i=1}^m (f_{\vec{w},b}(\vec{x}^{(i)}) - y^{(i)}) x_j^{(i)} \\ \frac{\partial}{\partial b} J(\vec{w}, b) &= \frac{1}{m} \sum_{i=1}^m (f_{\vec{w},b}(\vec{x}^{(i)}) - y^{(i)}) \end{aligned}$$

Then, replacing to the formulas of w_j and b :

$$\begin{aligned} w_j &= w_j - \alpha \left[\frac{1}{m} \sum_{i=1}^m (f_{\vec{w},b}(\vec{x}^i) - y^i) x_j^{(i)} \right] \\ b &= b - \alpha \left[\frac{1}{m} \sum_{i=1}^m (f_{\vec{w},b}(\vec{x}^i) - y^i) \right] \end{aligned}$$

Although the update rules resemble those of linear regression, the key difference is in the function $f_{\vec{w},b}(\vec{x})$. In logistic regression, it uses the **sigmoid** applied to $\vec{w} \cdot \vec{x} + b$, making it fundamentally different despite the similar form.

- Linear Regression: $f_{\vec{w},b}(\vec{x}) = \vec{w} \cdot \vec{x} + b$
- Logistic Regression: $f_{\vec{w},b}(\vec{x}) = g(\vec{w} \cdot \vec{x} + b) = \frac{1}{1 + e^{-(\vec{w} \cdot \vec{x} + b)}}$



Same concept:

- Monitoring Gradient Descent. (*learning curve*)
- Vectorized Implementation. (*to improve the efficiency of gradient descent*)
- Feature Scaling. (*enhance the convergence speed of gradient descent*)

Example:


Running Gradient Descent and the cost decreases for each loop.

```

Iteration 0: Cost 0.684610468560574
Iteration 1000: Cost 0.1590977666870456
Iteration 2000: Cost 0.08460064176930081
Iteration 3000: Cost 0.05705327279402531
Iteration 4000: Cost 0.042907594216820076
Iteration 5000: Cost 0.034338477298845684
Iteration 6000: Cost 0.028603798022120097
Iteration 7000: Cost 0.024501569608793
Iteration 8000: Cost 0.02142370332569295
Iteration 9000: Cost 0.019030137124109114

updated parameters: w:[5.28 5.08], b:-14.222409982019837

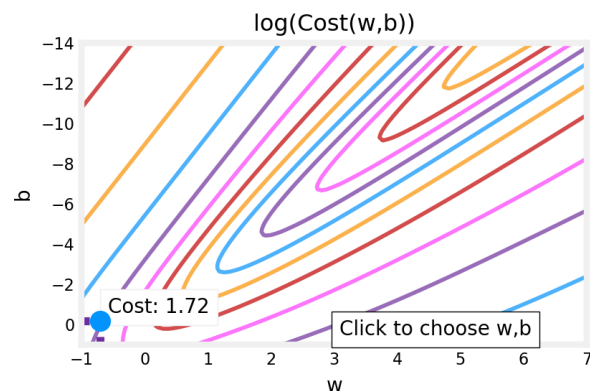

```




In the following plot:

- the shading reflects the probability $y = 1$ (result prior to decision boundary)
- the decision boundary is the line at which the probability $=0.5$


For this position of w and $b \Rightarrow$ the Sigmoid shape is not good



It is a little better



Better:



It is the best cost for logistic regression.

