

Additional Neural Network Concepts

Advanced Optimization

Adam optimization algorithm, which improves the training process of neural networks compared to traditional gradient descent.

Adam Algorithm Intuition

- **Adam: Adaptive Moment Estimation** and adjusts the learning rate automatically based on the behavior of the parameters during training.
- It uses different learning rates for each parameter, allowing for more efficient convergence to the minimum of the cost function.
- If w_j (or b) keeps moving in the same direction, **increase** α_j .
- If w_j (or b) keeps oscillating, **reduce** α_j .

MINIST Adam

```
model = Sequential([
    Dense(units=25, activation='relu'),
    Dense(units=15, activation='relu'),
    Dense(units=20, activation='linear'),
])
# it has some initial learning rate
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-3),
              loss=SparseCategoricalCrossEntropy(from_logits=True))

model.fit(X, Y, epochs=100)
```

- It is more robust to the choice of initial learning rate, making it a preferred option among practitioners for training neural networks.

Additional Layer Types

Recap: Dense Layer

Every neuron in the layer gets its inputs all the activations from the previous layer.


$$a_j^{[l]} = g(\vec{w}_j^{[l]} \cdot \vec{a}^{[l-1]} + b_j^{[l]})$$

Convolutional Layer

If **each neuron (unit)** in a layer only looks at **a subset (part)** of the **outputs** from the previous layer, this approach can offer certain advantages.



- Faster computations.
- Less training data (Less prone to overfitting).

The example illustrates how a convolutional layer processes a handwritten digit image by having each neuron look at a limited rectangular area of the image.



Convolutional Neural Network

In a different context, the convolutional layer can be applied to **EKG signals (electrocardiograms)**, where neurons analyze small windows of the signal to classify heart conditions. (it is 1D array not 2D array example)




- Each neuron in the convolutional layer examines a small window of the EKG signal:
 - For example, the first neuron might analyze the first 20 data points, while the second

neuron looks at points 11 to 30.

Second Convolutional Layer:

- The second layer takes the activations from the first layer as its input.
- Each neuron in this layer also looks at a limited number of activations from the previous layer, rather than all of them.
- For example:
 - The first neuron in the second layer might look at the first 5 activations from the first layer.
 - The second neuron might analyze activations 3 to 7.
 - The third neuron might focus on activations 5 to 9.
- This method allows the network to identify patterns in the EKG signal that can indicate heart conditions, improving classification accuracy.




Purpose of the Second Layer:

- This structure allows the network to learn more complex features by combining the information from the first layer's activations.
- By focusing on different subsets of activations, the second layer can capture higher-level patterns in the EKG signal that may be indicative of specific heart conditions.

Final Output Layer

- The final output layer (often a sigmoid layer) takes the activations from the last convolutional layer.
- This output layer typically makes a binary classification decision, such as determining whether the patient has a heart disease or not, based on the combined information from the previous layers.



QUES: The Adam optimizer is the recommended optimizer for finding the optimal parameters of the model. How do you use the Adam optimizer in TensorFlow?

- The call to `model.compile()` will automatically pick the best optimizer, whether it is gradient descent, Adam or something else. So there's no need to pick an optimizer manually.
- When calling `model.compile`, set `optimizer=tf.keras.optimizers.Adam(learning_rate=1e-3)`. 🍌🍌
- The Adam optimizer works only with Softmax outputs. So if a neural network has a Softmax output layer, TensorFlow will automatically pick the Adam optimizer.
- The call to `model.compile()` uses the Adam optimizer by default

Explain: Correct. Set the optimizer to Adam.

```
model = Sequential([
    Dense(units=25, activation='relu'),
```

```
Dense(units=15, activation='relu'),  
Dense(units=20, activation='linear'),  
    ])  
# it has some initial learning rate  
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-3),  
              loss=SparseCategoricalCrossEntropy(from_logits=True))  
  
model.fit(X, Y, epochs=100)
```